```
────────────────────────────────────────────────────────────────────────────────
MOMENTUM REVERSAL RISK — formation 2020-10-31 (last session 2020-10-30)
US equity cross-sectional momentum, the reconstructed 12-1 book:
top and bottom deciles of the Russell 3000, VALUE-WEIGHTED within each leg.
Horizon: 10 trading days at the 5% quantile. The probability line below is
    the VaR
restated at the event threshold, not a forecast of a reversal's timing.
────────────────────────────────────────────────────────────────────────────────

SEVERITY   [the parameter-free scaled quantile; no conditioning state]
  10-day 5% VaR      -19.29%     against an unconditional -7.57%   (2.5x)
  10-day 5% ES      -31.93%
  the book's trailing volatility is 60.9% annualised, and the estimator is
      today's
  volatility times the 5% quantile of past scaled returns. Nothing is fitted.
  P(10-day return at or below the book's own unconditional -7.57%): 20.6%.
  At French's long-run reference of -4.32%, which is a different and lower
      bar for a
  book this volatile, the same figure is 29.3%.

  Three conditioning variables were tested and none is used: the bear
      state, which
  made the forecast worse; the calendar state; and the volatility tercile.
      Their
  ratios and intervals are in the register.

IS THIS LINE CALIBRATED   [out of sample, on the book, formations BEFORE
    this date]
  scaled line          5 breaches of 54, 9.26%, Kupiec p 0.197  -> passes
  unconditional        9.26%, Kupiec p 0.197  -> passes
  At this date the two lines cannot be told apart on count: each breaches
      on the
  same number of formations, 4 of them the same formations, and Kupiec says
  the same of both. They are not the same set -- each has one breach the other
  does not -- so this is agreement on the test, not identical behaviour. The
  divergence that makes the scaled line the better calibrated one appears
      later
  in the tier, and a reader standing here could not know it.

WHAT THE BOOK IS   [value-weighted, as the design specifies]
  long leg            263 names, but an effective 8.0 of them; the largest
      is 27%
                     of the leg and the five largest 60%
  short leg          an effective 11.7 names
  This is what value-weighting inside a decile produces. It is a property
      of the
  construction and is printed, not corrected.

WHAT THE READER'S EXISTING SCREENS SAY
  Sector screen      263 names across 10 GICS sectors, Herfindahl 0.146
                     largest Energy at 23%  ->  reads DIVERSIFIED
  VIX                38.0, 96.8% of its own history  ->  high, no direction
```

WHAT THE LOSER LEG IS PRICED ON [8-K filings, gated on acceptance time]

The theme below is the 8-K survival-condition reading, by the fallback rule registered before either result was seen. The theme pipeline ran on this date and produced no theme above the placebo.

The three sentences that follow are INSTRUMENT DIAGNOSTICS, not market data, and they are NOT point-in-time: they score the pipeline across every date it was run on, five of which are after this one and three after the tier seal. They say how well the instrument works, which a reader here could not know; they are printed because the alternative is asking you to trust the fallback without saying why it was needed.

Across 8 episode dates it raised a theme on 2. The registration named an expected label in advance for 3 of those dates, and 0 matched -- but no match was reachable, because the dates where a theme rose are not among the dates that carry an expected label. That zero is a property of which dates were labelled, not evidence about the labels.

The reason is coverage, not the placebo: on the dates where something did rise, the excess weight rested on a handful of companies with Item 1A text out of a set of about a hundred and thirty. That is reported rather than omitted, and it is why the classifier run is listed as not built.

Stated an external survival condition: 19 of 112 filings (17.0%)

9 names carry one, and they name the same thing: the resumption of normal demand after COVID-19. The sector screen above sees ten sectors. The filings see one bet.

Evidence, verbatim, with the date it became public:

- UAL accepted 2020-09-09
  "expects demand to remain suppressed and plateau at levels of around 50%, relative to 2019 levels, until a widely accepted treatment and/or vaccine ..."
- CCL accepted 2020-08-07
  "As a result of the COVID-19 outbreak, we have paused our guest cruise operations, and if we are unable to re-commence normal operations in the near..."
- AMC accepted 2020-09-24
  "We currently estimate that unless theatre attendance levels improve significantly from the third quarter of 2020 to the fourth quarter of 2020 and ..."

WHEN DID THIS LAST HAPPEN   [live-only; not backtestable; run on a historical date]

Five nearest dates by the deterministic distance rule, and what followed each:

- #1 2009-08-31   distance 0.039   forward 21d -4.03%   conditions: NO RECORD
- #2 2009-05-31   distance 0.070   forward 21d +3.91%   conditions: NO RECORD
- #3 2009-01-31   distance 0.085   forward 21d +11.24%   conditions: NO RECORD
- #4 2008-10-31   distance 0.088   forward 21d +6.33%   conditions: NO RECORD
- #5 2000-05-31   distance 0.110   forward 21d +16.09%   conditions: NO RECORD

Four of the five were followed by gains. Retrieval alone would reassure.

A three-role panel argued over that record. 7 claims survived the gate and 1 was dropped: a mechanism claim about a date whose condition record could not support it. The gate keeps a mechanism claim only with a verbatim span from a condition record for that date, and counts what it drops.
Adjudicator: the match is NOT APT and mechanism CANNOT BE RULED ON, deciding on the field condition_record_status -- every one of the five predates the condition corpus, so what they were priced on is unknown rather than different.

Two limits, both real. This step is LIVE-ONLY: it is the design's one agentic component, a production system would use it, and it cannot be backtested, so its contribution here is shown and not measured. And the distance used 3 blocks of the v1 engine's state space, not the design's exposures, so these are the v1 engine's neighbours.

CATALYST DAYS IN THE WINDOW   [fields, not a model input]

An FOMC decision day falls in the next ten sessions. FOMC dates are published a year ahead, so this one is genuinely scheduled and genuinely knowable today.
The earnings figure this line used to carry has been REMOVED. It was built from the acceptance timestamps of 8-K item 2.02 filings inside the forecast window -- that is, from filings that did not exist on this date. It described what did happen, not what was scheduled, and a page that sells point-in-time discipline cannot print realised data under the word `scheduled`. Reporting an earnings calendar properly needs a source of announced dates, which this package does not have; it is named in DESIGNED AND NOT BUILT rather than approximated.
Neither field is an input to the severity line, because conditioning on the calendar state did not improve it.

WHAT WOULD CHANGE THIS READING

A vaccine or treatment readout resolves the shared condition, the loser leg reprices upward together, and the short half takes the loss. Named by the filings, not by a model: UAL states it as the condition on 2020-09-09.

DESIGNED AND NOT BUILT [named rather than given a number]

Positioning and crowding, the vulnerable-name lists, the catalyst agent, and the theme pipeline's classifier run. None appears above as a figure.

WHAT HAPPENED NEXT [NOT part of the page: outside the information set of 2020-10-30]

Realised over the ten sessions 2020-11-02 to 2020-11-13: -14.58%.
That is beyond the unconditional line and inside the volatility-scaled one.
2020-11-09 alone was -20.21%: the winner leg -4.18%, the loser leg +16.03%. The loss is the short half, which is the mechanism this page named in advance under WHAT WOULD CHANGE THIS READING -- a vaccine readout resolving the shared condition and repricing the loser leg upward together.
French's published factor over the same window: -9.56%. This book is the more concentrated portfolio and falls further.
One formation is one observation. It does not make the scaled line right; the coverage tests over 54 prior formations do that, and they are printed above.